

Advanced Systems Lab Report

Autumn Semester 2018

Name: David Yenicecik
Legi: 15-944-366

Grading

Section	Points
1	
2	
3	
4	
5	
6	
7	
Total	

Notes on writing the report (remove this page for submission)

Furthermore, it is expected that the interactive law holds for all experiments. In case throughput and response time do not match, it is imperative that you explain why, otherwise you risk losing most of the points for the experiment in question.

1 System Overview (75 pts)

I structure my project (code only) into the following folders. I give a short explanation for each item on how it is used.

I will start out with a diagram that explains the overall structure and refers to each .java file. I will then be more detailed with my project-code structure.

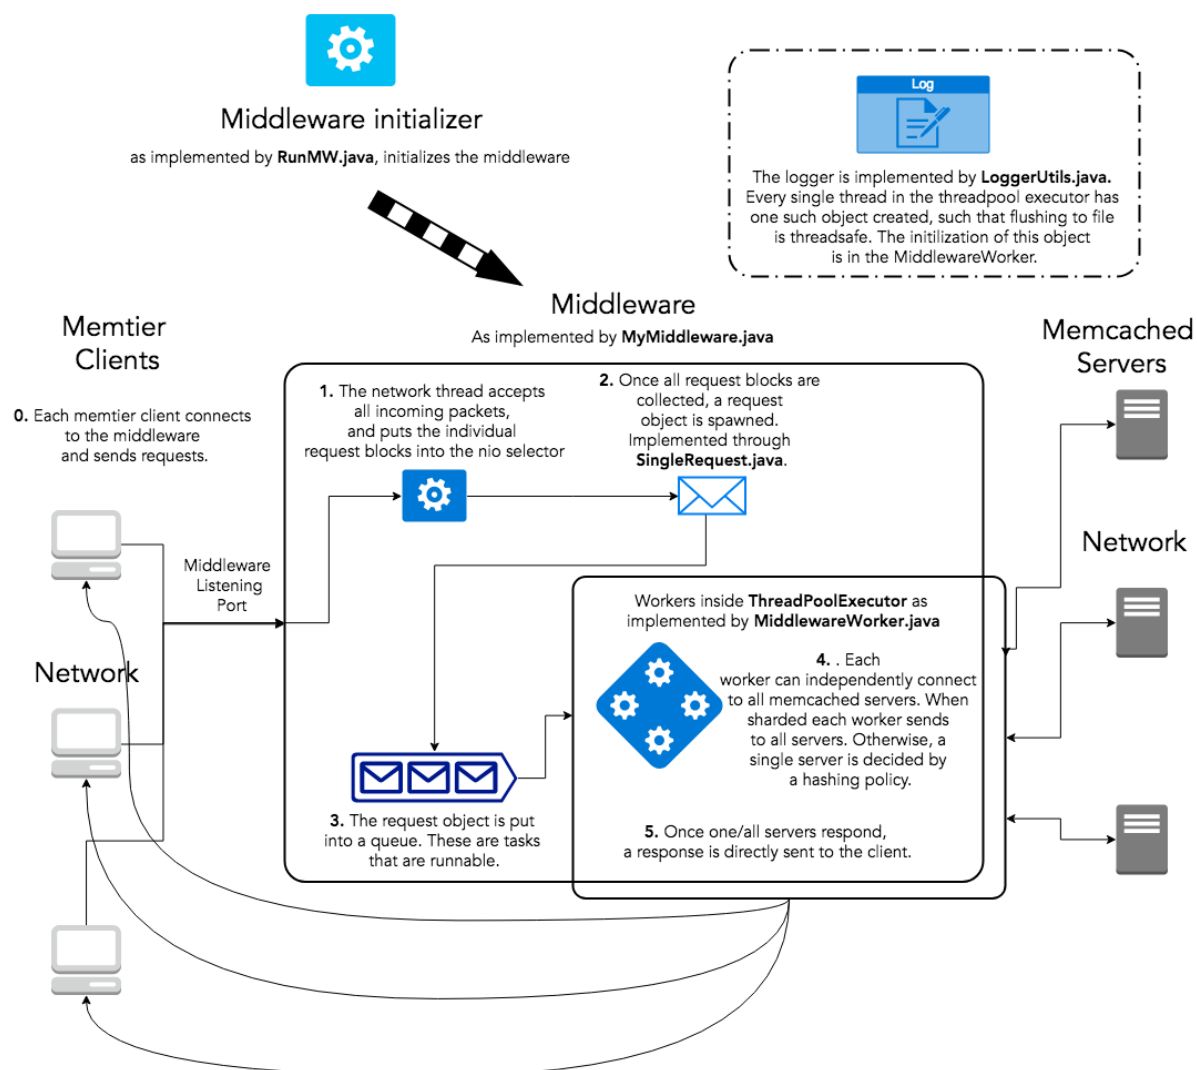


Figure 1: Exp.2.1: A figure with two subfigures

1. scripts

This folder includes all the logic to automatically compile the code, deploy the code, run the experiments (each individual one), and automatically download the logs. This is true both for a local development environment (using docker VM's), and the external

"production" environment (using Azure VM's). The local docker and external Azure systems are interchangeable with a small change in command.

2. **data**

This folder includes all the raw data that is pulled from the experiment, all the python files which process upon this data, and all the processed data.

3. **figs**

This folder includes all the figures that are create using the processed data from point 2. I use python to create figures from individual graphs

4. **src**

I will talk about the code structure of the src directory with more detail in this next section. I take a top-down approach while explaining (i.e. starting with where requests originate, and where they go from there). I will keep this concise, as the *lifetime of a request* section covers some information on what happens in which file.

(a) **RunMW.java**

This is the entry point / main class of the program. This is the default implementation by the TA's.

(b) **SingleRequest.java**

Is the class which encodes a single request from one of the clients. This can include the SET or GET operation. For logging purposes, this class also includes all possible times which may be of use to calculate the queueing theory (and latency and throughput) later on.

The request type is parsed by checking the very first character of the request string. If the string starts with a "g", the request type is a GET (later on, it is decided on the fly if it is a MULTIGET by checking if the number of keys is greater than 1). If the string starts with a "s", the request type is a SET. If none of the above cases hold, an error happened (we exit gracefully, as this is unintended behavior).

(c) **MyMiddleware.java**

This is the entryptpoint of how the Middleware is called. The datastructure we use to connect to the individual clients is the **nio.Selector**. The **nio.Selector** can hold multiple connections from different clients conneting. Furthermore, it can parse packages that don't immediately fit into the bytearray. The backlog size is bigger than 0 such that multiple requests can made to the same channel at the same time (and these requests are backlogged).

This class is responsible for fetching the individual requests (spawned as a *SingleRequest* object) using non-blocking IO, putting it to a queue, and spawning *MiddleWareWorker*'s within a **ThreadPool** to act upon these requests. We use the NIO selector to handle all connections. The entire logic of this .java class runs in on the main thread of the class, which I may refer to as "Network Thread". This allows for multiple clients to connect to the server in a non-blocking fashion. Whenever a connection is compelte (i.e. the request is complete), this selector spawns a *MiddleWareWorker* and a *SingleRequest*.

The diagram above symbolizes how this works.

(d) **MiddleWareWorker.java**

The *MiddleWareWorker* takes a *SingleRequest* and passes it to the server(s). The

MiddlewareWorker implements **Runnable**, and is executed by the **ThreadPoolExecutor** by calling **submit**. Depending on whether the request type is a SET or GET, we have different behavior:

- case SET: The *SingleRequest* is sent to each individual server in a sequential manner. After the sending to the server is done for all servers, it listens to the response of each individual server. It listens until all servers have responded. If any single item has responded with an error, this SET operation responds with the first error encountered to the client. Otherwise it returns a **STORED** message.
- case GET: The string of the *SingleRequest* is used to calculate a hash for the individual request. This hash is then modulo-ed with the number of server that we can send the request to. The modulo operation decides which server to send it to. This has **provenly uniformly at random distribution** behavior (proof at:). In short, this proof relies on the fact that hashing is pseudorandom, and that pseudorandomity means that for any input, all output values are uniformly at random distributed. Pseudorandomity keeps it's properties when applied with the modulo operator. Because all operations preserve pseudorandomity, the final operation - and thus the server chosen - is also uniformly at random.
- case MULTIGET (nonsharded): The nonsharded multiget acts exactly like the GET case. Again, the uniformly at random assumption is guaranteed because the hashing algorithm is pseudorandom (and as such provides keys uniformly at random).
- case MULTIGET (sharded): The sharded MULTIGET case acts as follows. The MULTIGET request is first split up into $n = \max(keys, buckets)$ where *buckets* is the number of servers in total, and *keys* is the number of keys in total. Each individual split up request is then treated as an individual GET request. When the memcached server responds, all the answers are concatenated in a sequential fashion and sent back to the client. Errors are handled and also reduced to the first occuring error if this is the case.

(e) **LoggerUtils.java**

A helper class. Includes all the logic that is needed to log the requests to hard disk. All the request logic is accumulated to variables, and the mean is flushed to disk every few minutes. For GET requests (and in the interest of experiment 5), GET requests are accumulated into a list, and flushed to disk every 5 seconds. The logging happens in such a way that each individual *MiddlewareWorker* has it's list of requests and all accumulator objects (inside the LoggerUtils.java) to be logged. Because there is a separate Logger for each thread, there are no issues with multithreading as there is no concurrent data access.

Every single request keeps track of the following values, which is then flushed to the LoggerUtils.java (and thus to the file) object when the request has been successfully. All the information in SingleRequest is used to create the following log-informations.

- i. timeRealOffset
- ii. differenceTimeCreatedAndEnqueued
- iii. differenceTimeEnqueuedAndDequeued
- iv. differenceTimeDequeuedAndSentToServer
- v. differenceTimeSentToServerAndReceivedResponseFromServer

- vi. `differenceTimeReceivedResponseFromServerAndSentToClient`
- vii. `timeRealDoneOffset`

(f) **RequestType.java**

A helper struct definition, which defines the two possible input types (Multi-gets are decided on the fly at a different point as described in *SingleRequest.java*)

1.1 Lifetime of a request

In the following I will talk about how requests enter the middleware, how they are parsed, how they get distributed to servers, and how the middleware communicates these values back to the clients. This is a more detailed version of the above diagram.

1. Request coming from client to middleware:

When a request comes from a client to the middleware, I use an **nio.channels.Selector** to accept the request. This datastructure has the following benefits. First, it can distinguish between multiple clients. Second, it can process these individual requests simultaneously in an asynchronous (non-blocking) manner. Third, it fills up multiple channels (one for each connection, and thus, one for each client), which means that if the request does not fit into one network packet, it will just listen for the rest of the packet. I detect if a single request fills up by filling a java **ByteBuffer** until it the request has come to an end (which we can recognize by waiting for the **END** keyword. We parse the type of request by looking at the very first character (interpreting the bytes) and cross-comparing if this is a *set* or *get*. To distinguish between *multi-get* and *gets*, we later on split the string resulting from parsing the bytebuffer by spaces. If the number of elements after splitting is bigger than 2 (the "get" keyword, and the key), then I parse a multi-get. Else, I parse a get.

2. The incoming request spawns a SingleRequest object:

The SingleRequest object is specified in the **SingleRequest.java** java file and is wrapped around a **MiddlewareWorker.java** object which implements a java **Runnable** on which I later on call **submit** using a **ThreadPoolExecutor**. This file keeps track of the statistics described in the **LoggerUtils.java** class for logging. The SingleRequest takes over the **ByteBuffer** which was created while **nio.channels.Selector** was listening for a complete network packet. This ByteBuffer will later on be passed to the individual server(s) (after some processing).

3. Submitting a SingleRequest to a MiddlewareWorker using the java ThreadPoolExecutor:

I use the SingleRequest that was generate before, and spawn a new **MiddlewareWorker**. This **MiddlewareWorker** can then be submitted to the java **ThreadPoolExecutor** which contains the number of middleware-threads (as specified per experiment). Each individual middlewareworker contains one instantiated **LoggerUtils** class per thread and thus is threadsafe.

4. Sending the server response back to the client:

Each middlewareworker, and thus each individual thread, has a connection to the client that the request came from, and a hashmap of sockets to connect to all the servers. The response is sent back to the client using the description in the previous sub-section. In any case, this goes through the individual middleware-worker thread, and does **not** go through the initial **nio.channel.Selector** again.

1.2 Some statistics, and observed bandwidths amongst different VMs

Some general observations include the following approximate maximum throughput. I did not include graphs to this, as all these experiments just involve scalars.

I use the program *iperf -c 'VM Address'* to arrive at these statistics. I arrive at the theoretical maximum throughput by checking how many packages could possibly fit in the bandwidth, where the message size is 4KB, and the bandwidth is specified by the value before that. I divide the bandwidth by the packet size to arrive at the theoretical maximum throughput per single VM.

VM From	VM To	Bandwidth	Theoretical Maximum Throughput per VM (ops/sec)
Client	Middleware	201 Mbits/sec	6'250
Client	Server	201 Mbits/sec	6'250
Middleware	Client	804 Mbits/sec	25'000
Middleware	Server	804 Mbits/sec	25'000
Server	Client	101 Mbits/sec	3'125
Server	Middleware	101 Mbits/sec	3'125

Intuitively, this means that if we have 3 clients and one server, the server must be a bottleneck, and that the clients share the 3'125ops/s (approximately 1'000ops/s per client). **I will use these numbers to derive some explanations later on.**

In addition to the vanilla experiments, I am now talking about my statistical methods of how I analyse the experiments. Each plotted value for each experiment in the subsequent report was generating using **three repetitions**, each including a **warm-up**, and a **cool-down** phase of 15 seconds (where the core-length of the experiment was 60 seconds. To arrive at a single rate amongst rates, I use the **arithmetic mean**. As an **error mesaure**, I use the standard deviation. These are common statistical tools, and according to the **law of large numbers**, will approximate any distribution truthfully when the number of samples go against infinity, which is why I chose these two measures.

The mean μ and standard deviation σ are calculated as follows:

$$\mu = \frac{\sum_{i=1}^n x_i}{n} \quad (1)$$

$$\sigma = \sqrt{\frac{\sum_{i=1}^n (x_i - \mu)^2}{n}} \quad (2)$$

I will now proceed with the experiments.

2 Baseline without Middleware (75 pts)

These experiments don't use the middleware, and only consist of memtier client instances, and memcached server instances.

2.1 One Server

I use the following setup:

Number of servers	1
Number of client machines	3
Instances of memtier per machine	1
Threads per memtier instance	2
Virtual clients per thread	[1..32]
Workload	Write-only and Read-only
Repetitions	3 or more

I will test out the response time and latency for a different number of virtual clients per thread on the client-side (per memtier instance), namely [1, 2, 4, 8, 16, 32]. I will have two different subsections for read-only operations, and write-only operations. As there is no difference in pre-populating the servers in the case of writes, for both experiments I pre-populate the server the same way as in experiment 2.1 (Baseline without middleware and 3 clients) using the following command, which sequentially generates and stores each key in the server. **Whenever in future experiment setups I refer to prepopulating the server, I will refer to this command:**

```
memtier_benchmark -s {SERVER_IP} -p {PORT}
--protocol=memcache_text --clients={VIRTUAL_CLIENTS_PER_THREAD} --threads={THREADS}
--requests=15000 --ratio=1:0 --data-size=4096
--expiry-range=9999-10000 --key-maximum=10000 --key-pattern=S:S
```

The clients and servers don't use multi-gets, and there is no middleware involved. I run 3 repetitions of each configuration, each having a length of 90 seconds (such that a warm-up and cool-down time of 15 seconds respectively evens out).

I run 3 repetitions of each configuration, each having a length of 90 seconds (such that the warm-up and cool-down times even out).

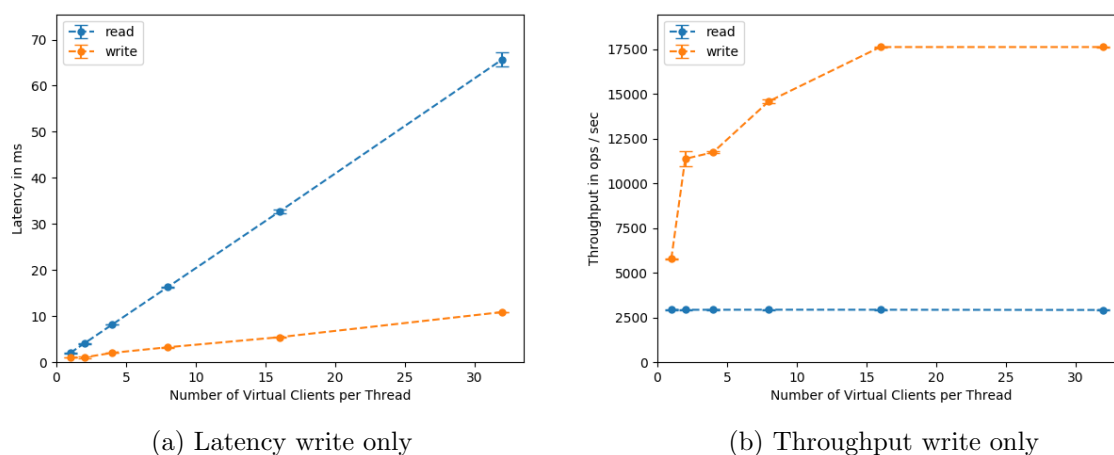



Figure 2: Exp.2.1: A figure with two subfigures

2.1.1 Read-only

I first present the graph. The throughput reached approximately almost $3000ops/s$, which conforms to the observation of the maximum throughput per server in section 1. Please notice that the error is very small, and thus for some points intersects with the measured points. The error metric is explained in section 1.



(a) Response time / Latency in ms per single request

(b) Total throughput pre

Figure 3: Experiment 2.1, Baseline without middleware and one server. Read-only latency and throughput per individual client (mean \pm standard deviation)

For read-only operations, the bottleneck is the upload bandwidth of the server, as 3 clients are trying to download a load of 100Mbit/s, and thus must each share appr. 33Mbit/s. This corresponds to the total upload bandwidth of $3000ops/sec$ (which was empirically proven in section 1 through an additional experiment), is thus divided amongst three servers. The missing ops/sec (between our observers ops/sec and 3000 comes from network overhead. This also proves to be a valuable sanity check. The clients are able to generate enough load with even 1 virtual thread per thread, such that with 1 virtual client per thread the network bandwidth is saturated (and stays like this for an increasing number of virtual clients per thread). As the number of threads increases, more requests are able to be generated. Because the network bandwidth stays constant, but we introduce more virtual clients, the round-trip time of individual requests increases. This implies a linear increase in latency, as can be seen from the graph. This linear increase in the response time provides a second sanity check

2.1.2 Write-only

Again, I first present the graph. The throughput reached approximately almost $18000ops/s$, which conforms to the observation of the maximum throughput per server in section 1. Please notice that the error is again very small, and thus barely readable around the measured points. The error metric is explained in section 1.

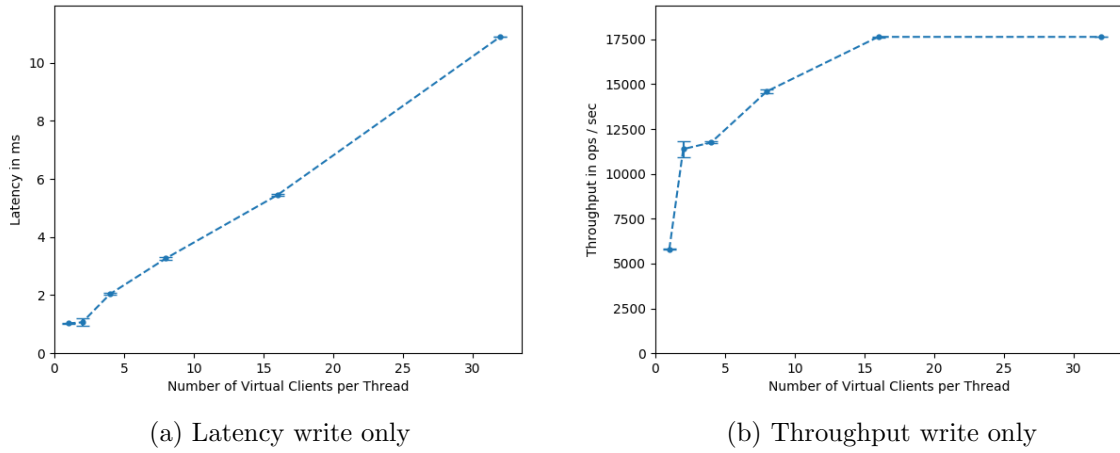


Figure 4: Exp.2.1: A figure with two subfigures

For write-only operations, the bottleneck is the upload bandwidth of the three clients, as 3 clients are trying to upload a load of 200Mbit/s each. The server only responds with message suchs *STORED*, which take up almost no bandwidth compared to the actual message itself. Thus, the total bandwidth is the additive bandwidth of each individual client, which means approximately $3 \times 200Mbit/s = 600Mbit/s$, which corresponds to approximately 18000ops/sec (same calculation as in section 1.2). The missing ops/sec (between our observers ops/sec and 18000 comes from network overhead. This also proves to be a valuable sanity check. I know that storing the individual requests is not the bottleneck, as a local docker experiment proves that higher throughputs can be achieved locally. The clients are not able to generate enough load with 1, 2 or 4 virtual clients per thread, and only generate a maximum load with 16 virtual clients per threads, which saturates the upload bandwidth of the client VMs. This can be seen as the throughput grows square-root like, and then saturates. This can also be seen from the graph, which plateaus around 16 virtual clients per thread. As a sanity check, the latency graph shows how the latency does not increase for adding additional virtual clients per thread (although this should increase linearly). Only after the 4th virtual client per thread, does the response time increase linearly, which underlines the saturation phase.

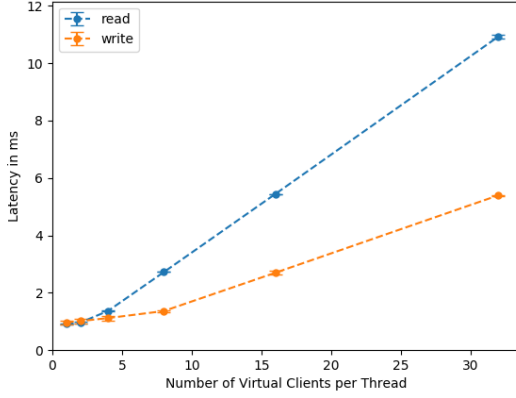
2.2 Two Servers

For the following experiment, I use the following setup:


Number of servers	2
Number of client machines	1
Instances of memtier per machine	2
Threads per memtier instance	1
Virtual clients per thread	[1..32]
Workload	Write-only and Read-only
Repetitions	3 or more (at least 1 minute each)

I will test out the response time and latency for a different number of virtual clients per thread on the client-side (per memtier instance), namely [1, 2, 4, 8, 16, 32]. I will have two different subsections for read-only operations, and write-only operations. As there is no difference in pre-populating the servers in the case of writes, for both experiments I pre-populate the server the same way as in experiment 2.1 (Baseline without middleware and 3 clients). The clients and servers don't use multi-gets, and there is no middleware involved. I run 3 repetitions

of each configuration, each having a length of 90 seconds (such that the warm-up and cool-down times even out).



(a) Latency write only




(b) Throughput write only

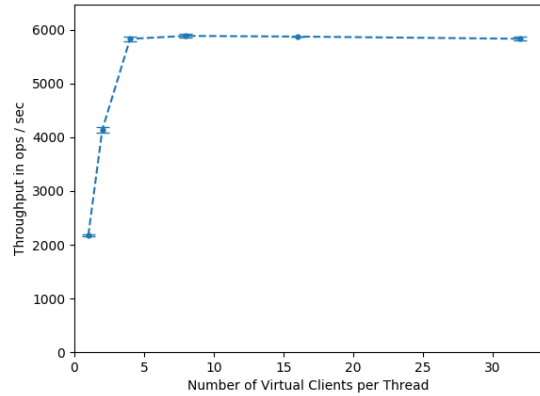
Figure 5: Exp.2.1: A figure with two subfigures

2.2.1 Read-only

I first present the graph. The throughput reached approximately almost $6000ops/s$, which conforms to the observation of the maximum throughput per server machine in section 1. This means that both server virtual machine, are using their maximum capacity to upload all the values that are uploaded by the server virtual machines. The combined throughput of the two servers is $2 \times 100Mbit/s$ which corresponds to a compound $200Mbit/s$. This conform to the $6000ops/sec$ when using the same calculation as in section 1.2. Please notice that the error is again very small, and thus intersects with the measured points. The error metric is explained in section 1.



(a) Exp2.2: Latency read only



(b) Throughput read only

Figure 6: Exp2.2: A figure with two subfigures

For read-only operations, the bottleneck is the upload bandwidth of the server virtual machines, as each server is uploading $100Mbit/s$ each, which compounds to $200Mbit/s$ together. The total upload bandwidth of $6000ops/sec$ (which was empirically proven in section 1 through

an additional experiment), is thus divided amongst amongst two memtier servers and thus two memtier instances on the same machine. This also proves to be a valuable sanity check.

The server can respond fast enough to all requests by the clients, as long as there is only up to 4 virtual clients per thread. Once the client spawns 4 virtual clients per thread, the server VM bandwidth starts to saturate. The system does not become stable but saturated after an increase of virtual client threads, as can be seen from the plot which indicates that there is no oversaturation. Because the network bandwidth stays constant, the round-trip time of individual requests increases, except for up to 4 virtual clients per thread. For 4 virtual clients per threads, the response time stays constant, simply due to the fact that the bandwidth is not saturated, and the requests don't have to wait on each other to be processed (but can be processed in parallel). This implies an almost constant latency for up to 4 virtual clients per thread, and then an increase in latency, as can be seen from the graph. The latency and throughput graphs show a correct, inverse correlation. This almost constant, and linear increase in the response time as seen on the graph provides a second sanity check.

2.2.2 Write-only

Again, I first present the graph. The throughput reached approximately almost $6000ops/s$, which conforms to the observation of the maximum throughput per client machine in section 1. This means that both client virtual machine, are using their maximum capacity to upload all the values to the server virtual machines. The combined throughput of the two servers is $2 \times 100Mbit/s$ which corresponds to a compound $200Mbit/s$. This conform to the $6000ops/sec$ when using the same calculation as in section 1.2. Please notice that the error is again very small, and thus intersects with the measured points. The error metric is explained in section 1.

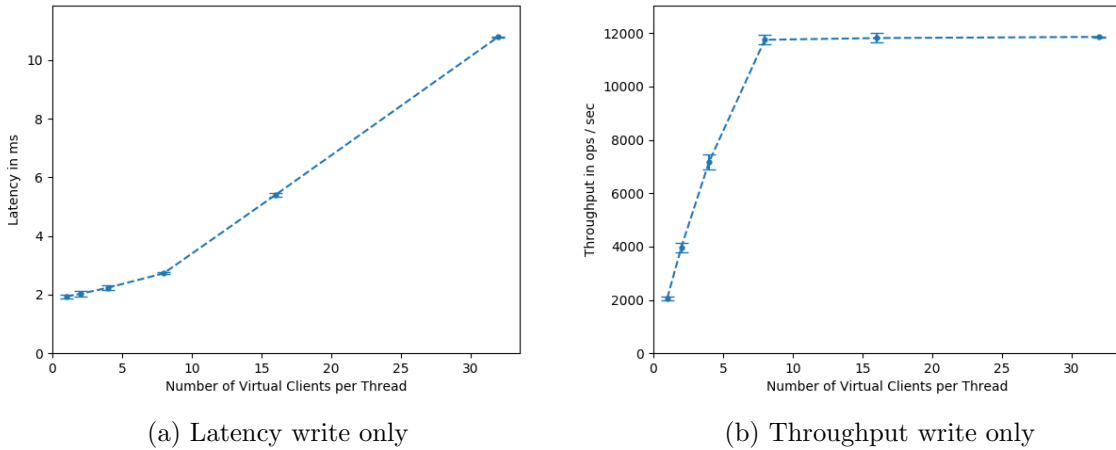


Figure 7: A figure with two subfigures

For write-only operations, the bottleneck is the upload bandwidth of the client virtual machines, as the single client is uploading $200Mbit/s$ each, which compounds to $200Mbit/s$ together. The sane expectation would be a throughput of $6'000 ops/sec$ (as given per the bandwidth of the client virtual machine). However, notice that the memtier client json files return values of $6'000 ops/sec$ each. However, the throughput of this operation is at $12'000 ops/sec$ as measure by the client, which does not immediately make sense, and is an issue that is very difficult to debug.

After careful analysis, I arrive at the conclusion that this may be due to one of the following two factors. 1. When I replicate the memtier-instances, the confirmation is sent from the server to both memtier-instances (which would mean that the throughput is doubled). I run one additional test to check what would happen when we have a single server. The total throughput of the system halves (as we have halved the servers). This indicates that the replicated confirmations may indeed be an issue causing the higher-than-expected throughput. This is unfortunately one factor that I cannot solve without changing the program of memtier, or modifying the virtual-machine setups. 2. Microsoft Azure may be responsible for some internal handling of the virtual-machine resource allocations, allowing for a higher bandwidth. Unfortunately, I cannot analyse this effectively as I do not have access to Microsoft Azure's policy of resource allocation, so explanation 1. is the only good explanation for this unexpected behavior.

The server can respond fast enough to all requests by the clients, as long as there is only up to 8 virtual clients per thread, because the response message (which is usually "STORED") consumes almost none of the bandwidth. Once the client spawns 8 virtual clients per thread, the server VM bandwidth starts to saturate as the client upload bandwidth of 200Mbit/s is reached. The system becomes saturated after an increase of virtual client threads, as can be seen from the plot which indicates that there is no oversaturation. Because the network bandwidth stays constant, the round-trip time of individual requests increases, except for up to 4 virtual clients per thread. For 4 virtual clients per threads, the response time stays constant, simply due to the fact that the bandwidth is not saturated, and the requests don't have to wait on each other to be processed (but can be processed in parallel). This implies an almost constant latency for up to 4 virtual clients per thread, and then an increase in latency, as can be seen from the graph. The latency and throughput graphs show a correct, inverse correlation. This almost constant, and linear increase in the response time as seen on the graph provides a second sanity check.

Latency and throughput correlate in the following way. The latency is inversely proportional to the throughput. One can see clearly that this explanation is supported by the two graphs, which show that overall trend of the latency is exponential, while the overall trend of the throughput is logarithmic. These are inverse functions of each other, as such provide a sanity check that the experiments go as expected.

2.2.3 Explanation

TODO: compare two-servers and one-server approaches.

Describe how this experiment compares to the previous section. Which results are the same and which ones differ? Explain what further conclusions can be drawn from the experiment.

2.3 Summary

Based on the experiments above, fill out the following table:

Maximum throughput of different VMs.

	Read-only workload	Write-only workload	Configuration gives max. throughput
One memcached server	1 / 2938	16 / 14595	VC=16
One load generating VM	4 / 5872	8 / 11859	VC=8

We can compare the read-only and write-only workloads. The write-only workloads have higher throughputs compared to the read-only workloads. Also, the configuration with only a

single memcached server has higher spread, which means that $17500 - 2900 = 14600$ is higher than $12000 - 5800 = 6200$. This is because we change it from one server to two servers, whereas we reduce the number of clients from 3 to 1. The change of number of clients has higher effect. In addition to that, saturation is reached much faster with only a single load-generating client, as all get and set request bodies (i.e. the data) has to enter and leave this VM. This means that in the load-generating VM, the single memtier client is the bottleneck. In the case of one memcached server, the bottleneck is the server (read only case), which has to respond with the response bodies and thus the maximum upload bandwidth (bottleneck) is reached. In the case of one memcached server write-only case, the bottleneck is the clients which have to send the request bodies. Each client has a maximum bandwidth of 6000 ops/sec.

3 Baseline with Middleware (90 pts)

In this set of experiments, I use three client memtier virtual machines, and 1 memcached server. These virtual machine instances are connected with exactly one middleware virtual machine in the middle. The three clients connect to the middleware. The middleware connects to the server.

The general tendency is that read-only requests are just as fast, while write-only requests are a bit slower and occur over-saturation in the middleware (probably due to the buffer sizes).

3.1 One Middleware

Number of servers	1
Number of client machines	3
Instances of memtier per machine	1
Threads per memtier instance	2
Virtual clients per thread	[1..32]
Workload	Write-only and Read-only
Number of middlewares	1
Worker threads per middleware	[8..64]
Repetitions	3 or more (at least 1 minute each)

The setup is exactly the same as in experiment "Baseline without Middleware and 1 server", with the difference that we inject one middleware between the clients, and the server.

In addition the measuring the throughput and response time for different values of number of virtual clients, we also allow to modify the number of middleware threads as another measurable variable. This means that I test out the throughput and latency for any permutation of $\text{virtualthreads}=[1, 2, 4, 8, 16, 32]$ and $\text{threads in the middleware}=[8, 16, 32, 64]$. I repeat each experiment for 3 times and plot the standard deviation amongst those trials. I also allow for a 15 second warm-up and 15 second cool-down time, and disregard these measurements when retrieving the logs about the request times from the middleware.

The throughput reached approximately almost 3000ops/s , which conforms to the observation of the maximum throughput per server in section "Baseline without Middleware and 1 server".

3.1.1 Read-only and Explanation

I will first plot the latency (response time) and the throughput as measured on the middleware. The throughput reached approximately almost 3000ops/s , which conforms to the observation of the maximum throughput per server in section 1. Please notice that the error is again very

small, and thus barely readable around the measured points. The error metric is explained in section 1. I will later show some plots generated by the logs of the client-machine which underline a correct measurement.

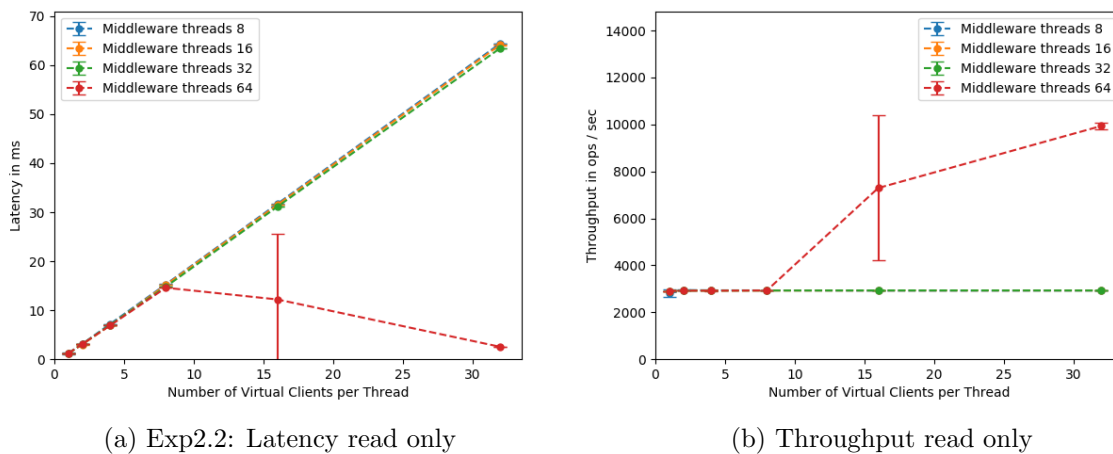



Figure 8: Exp3.1: Latency and throughputs as measures by the **middlewares**


For read-only operations, the bottleneck is the upload bandwidth of the client, as 3 clients are trying to download a load of 200Mbit/s, and thus must each share appr. 66Mbit/s. The client is a bottleneck before the middleware, which means that the middleware can easily manage all the read operations from the clients. The middleware virtual machine does not slow this down, as all GET requests are very short simple requests that all fit into the buffer. Also, the IO that serves as input to the middleware is non-blocking, which means that it can accept many such small requests simultaneously. This can be seen as increasing the number of middleware threads does not affect performance, as the client-side is bottlenecked. Each individual request is sent back to the client through it's designed thread, which does not block the response operation due to a large datasize. The total upload bandwidth of 3000ops/sec (which was empirically proven in section 1 through an additional *iperf* experiment), is thus divided amongst three clients. This also proves to be a valuable sanity check. The difference in the ops/sec (difference between our observers ops/sec and 3'125ops/sec is a reasonable network overhead.

The clients are able to generate enough load with even 1 virtual client per thread, such that with 1 virtual thread the network bandwidth is saturated. As the number of threads increases, more requests are able to be generated. Because the network bandwidth stays constant, the round-trip time of individual requests increases. This implies an increase in latency. The linear increase in the latency graph supports this.

As a sanity check, I present the throughput and latency plots from the client machines, which all conform to the throughputs and latencies as caluclated in the middleware. Another sanity check is that the interactive law holds, which states that the response-time **per request** is related inversely-linearly proportional to the throughput.



(a) Exp2.2: Latency read only




(b) Throughput read only


Figure 9: Exp3.1: Latency and throughputs as measures by the **clients**

The client and middleware graphs deviate slightly. As one can see, the middleware graph seems more stable. This is because the middleware graph only shows the performance **excluding** the warm-up and cool-down phase (each of which take approximately 15 seconds), whereas the client graphs do not exclude these measurements.

3.1.2 Write-only and Explanation



(a) Latency write only



(b) Throughput write only


Figure 10: Exp3.1: Latency and throughputs as measures by the middlewares

The single server is fast enough to respond to all individual client requests, as the server is only responding with the message "STORED" each time. However, at some point, the number of requests is so high, that either the individual client output bandwidth, or storing the actual requests on the hardware side becomes the bottleneck. This can be seen as the throughput grows square-root like, and then saturates. The throughput graph, which means out with 16 virtual clients per thread, proves this point. To decide whether the network is the bandwidth is the bottleneck, I repeat this experiment locally with virtual docker containers. These plateau at a later stage, which proves that the maximum network bandwidth is reached with 16 virtual clients per thread.


Latency and throughput correlate in the following way: ???

The following graphs are created using the logs of the client for the throughput and the latency, and underline that the values are sane.

As a sanity check, I present the throughput and latency plots from the client machines, which all conform to the throughputs and latencies as calculated in the middleware. Another sanity check is that the interactive law holds, which states that the response-time **per request** is related inversely-linearly proportional to the throughput.



(a) Latency write only



(b) Throughput write only

Figure 11: Exp3.1: Latency and throughputs as measures by the clients

The client and middleware graphs deviate slightly. As one can see, the middleware graph seems more stable. This is because the middleware graph only shows the performance **excluding** the warm-up and cool-down phase, whereas the client graphs do not exclude these measurements.

3.1.3 Additional Explanation

Provide a detailed analysis of the results (e.g., bottleneck analysis, component utilizations, average queue lengths, system saturation). Add any additional figures and experiments that help you illustrate your point and support your claims.

3.2 Two Middlewares

Number of servers	1
Number of client machines	3
Instances of mentier per machine	2
Threads per mentier instance	1
Virtual clients per thread	[1..32]
Workload	Write-only and Read-only
Number of middlewares	2
Worker threads per middleware	[8..64]
Repetitions	3 or more (at least 1 minute each)

The setup is exactly the same as in experiment "Baseline without Middleware and 1 server", with the difference that we inject two middlewares between the clients, and the server. Another difference is that we allow each individual client to run two mentier instances (each with one thread) to be able to connect to two instances each. The throughput reached approximately almost 3000ops/s, which conforms to the observation of the maximum throughput per server

in section "Baseline without Middleware and 1 server".


In addition the measuring the throughput and response time for different values of number of virtual clients, we also allow to modify the number of middleware threads as another measurable variable. This means that I test out the throughput and latency for any permutation of $\text{virtualthreads}=[1, 2, 4, 8, 16, 32]$ and threads in the middleware= $[8, 16, 32, 64]$. I repeat each experiment for 3 times and plot the standard deviation amongst those trials. I also allow for a 15 second warm-up and 15 second cool-down time, and disregard these measurements when retrieving the logs about the request times from the middleware.

Connect three load generator machines (two instances of memtier with $\text{CT}=1$) to two middlewares and use 1 memcached server. Run a read-only and a write-only workload with increasing number of clients (between 2 and 64) and measure response time *both at the client and at the middleware*, and plot the throughput and response time as measured in the middleware.


Repeat this experiment for different number of worker threads inside the middleware: 8, 16, 32, 64.

Number of servers	1
Number of client machines	3
Instances of memtier per machine	2
Threads per memtier instance	1
Virtual clients per thread	[1..32]
Workload	Write-only and Read-only
Multi-Get behavior	N/A
Multi-Get size	N/A
Number of middlewares	2
Worker threads per middleware	[8..64]
Repetitions	3 or more (at least 1 minute each)

3.2.1 Read-only




(a) Latency read only




(b) Throughput read only

Figure 12: Exp3.2: Latency and throughputs as measures by the middlewares



(a) Latency read only




(b) Throughput read only


Figure 13: Exp3.2: Latency and throughputs as measures by the clients

The client and middleware graphs deviate slightly. As one can see, the middleware graph seems more stable. This is because the middleware graph only shows the performance **excluding** the warm-up and cool-down phase, whereas the client graphs do not exclude these measurements.

3.2.2 Write-only

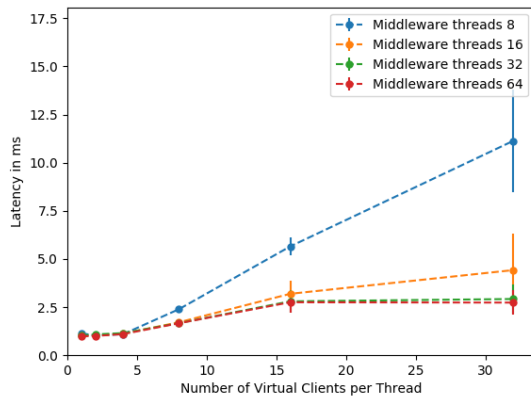


(a) Latency write only




(b) Throughput write only

Figure 14: Exp3.2: Latency and throughputs as measures by the clients



(a) Latency write only



(b) Throughput write only

Figure 15: Exp3.2: Latency and throughputs as measures by the middlewares

The client and middleware graphs deviate slightly. As one can see, the middleware graph seems more stable. This is because the middleware graph only shows the performance **excluding** the warm-up and cool-down phase, whereas the client graphs do not exclude these measurements.

3.2.3 Explanation

Provide a detailed analysis of the results (e.g., bottleneck analysis, component utilizations, average queue lengths, system saturation). Add any additional figures and experiments that help you illustrate your point and support your claims.

3.3 Summary

Based on the experiments above, fill out the following table. For both of them use the numbers from a single experiment to fill out all lines. Miss rate represents the percentage of GET requests that return no data. Time in the queue refers to the time spent in the queue between the net-thread and the worker threads.

Maximum throughput for one middleware.

	Throughput	Response time	Average time in queue	Miss rate
Reads: Measured on middleware	2933	64		0
Reads: Measured on clients	2987	64	n/a	0
Writes: Measured on middleware	8291	7		n/a
Writes: Measured on clients	7423	34	n/a	n/a

Maximum throughput for two middlewares.

	Throughput	Response time	Average time in queue	Miss rate
Reads: Measured on middleware	10330	31.40		0
Reads: Measured on clients	11116	32.37	n/a	0
Writes: Measured on middleware	8131	11.13		n/a
Writes: Measured on clients	8108	14.27	n/a	n/a

Notice that the miss rate is always zero, because this is 1. a closed system, and all servers are pre-populated before any experiment starts.

Based on the data provided in these tables, write at least two paragraphs summarizing your findings about the performance of the middleware in the baseline experiments.

4 Throughput for Writes (90 pts)

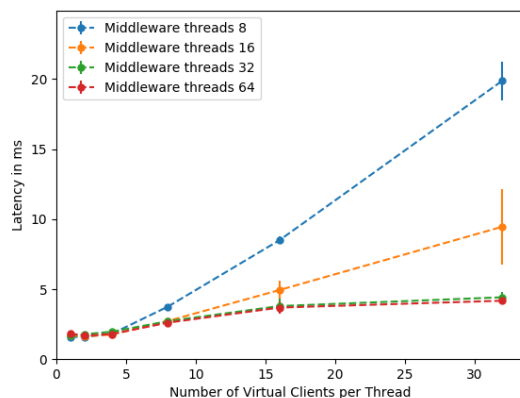
4.1 Full System

I am connecting three load generating client VMs to two middlewares. These middlewares are connected to three memcached server VMs each. I have the following setup.


Number of servers	3
Number of client machines	3
Instances of memtier per machine	2
Threads per memtier instance	1
Virtual clients per thread	[1..32]
Workload	Write-only
Number of middlewares	2
Worker threads per middleware	[8..64]
Repetitions	3 or more (at least 1 minute each)

During this experiment, I iterate over all possible permutations of virtual clients per thread (in the range of [1, 2, 4, 8, 16, 32]), and worker threads per middlewares (in the range of [8, 16, 32, 64]). I run each experiment for 90 seconds (which includes a 15 second warm-up and a 15 second cool-down time). Each experiment again consists of three trials from which we measure the mean and the standard deviation. This section only covers write-only experiments. I cover response time (latency) , and throughput.

The following are graphs from the middleware.




(a) Latency write only




(b) Throughput write only

Figure 16: Exp3.2: Latency and throughputs as measures by the middlewares



(a) Latency write only



(b) Throughput write only

Figure 17: Exp3.2: Latency and throughputs as measures by the clients

4.1.1 Explanation

Provide a detailed analysis of the results (e.g., bottleneck analysis, component utilizations, average queue lengths, system saturation). Add any additional figures and experiments that help you illustrate your point and support your claims.

4.2 Summary

Based on the experiments above, fill out the following table with the data corresponding to the maximum throughput point for all four worker-thread scenarios.

Maximum throughput for the full system

	WT=8	WT=16	WT=32	WT=64
Throughput (Middleware)	14365	18338	18828	17984
Throughput (Derived from MW response time)	14632	18827	18504	17830
Throughput (Client)	14858	17017	15703	17503
Average time in queue				
Average length of queue				
Average time waiting for memcached				

Based on the data provided in these tables, draw conclusions on the state of your system for a variable number of worker threads.

5 Gets and Multi-gets (90 pts)

For this set of experiments you will use three load generating machines, two middlewares and three memcached servers. Each memtier instance should have 2 virtual clients in total and the number of middleware worker threads is 64, or the one that provides the highest throughput in your system (whichever number of threads is smaller).

For multi-GET workloads, use the `--ratio` parameter to specify the exact ratio between SETs and GETs. You will have to measure response time on the client as a function of multi-get size, with and without sharding on the middlewares.

5.1 Sharded Case

Run multi-gets with 1, 3, 6 and 9 keys (memtier configuration) with sharding enabled (multi-gets are broken up into smaller multi-gets and spread across servers). Plot average response time as measured on the client, as well as the 25th, 50th, 75th, 90th and 99th percentiles.

Number of servers	3
Number of client machines	3
Instances of memtier per machine	2
Threads per memtier instance	1
Virtual clients per thread	2
Workload	ratio=1:<Multi-Get size>
Multi-Get behavior	Sharded
Multi-Get size	[1..9]
Number of middlewares	2
Worker threads per middleware	max. throughput config.
Repetitions	3 or more (at least 1 minute each)






Figure 18: Exp3.2: Latency and throughputs as measures by the middlewares




(a) Latency write only




(b) Throughput write only

Figure 19: Exp3.2: Latency and throughputs as measures by the clients



(a) Latency write only



(b) Throughput write only

Figure 20: Exp3.2: Latency and throughputs as measures by the middlewares

5.1.1 Explanation

Provide a detailed analysis of the results (e.g., bottleneck analysis, component utilizations, average queue lengths, system saturation). Add any additional figures and experiments that help you illustrate your point and support your claims.

5.2 Non-sharded Case

Run multi-gets with 1, 3, 6 and 9 keys (memtier configuration) with sharding disabled. Plot average response time as measured on the client, as well as the 25th, 50th, 75th, 90th and 99th percentiles.






Figure 21: Exp3.2: Latency and throughputs as measures by the middlewares

Number of servers	3
Number of client machines	3
Instances of memtier per machine	2
Threads per memtier instance	1
Virtual clients per thread	2
Workload	ratio=1:<Multi-Get size>
Multi-Get behavior	Non-Sharded
Multi-Get size	[1..9]
Number of middlewares	2
Worker threads per middleware	max. throughput config.
Repetitions	3 or more (at least 1 minute each)




(a) Latency write only




(b) Throughput write only

Figure 22: Exp3.2: Latency and throughputs as measures by the clients



(a) Latency write only



(b) Throughput write only

Figure 23: Exp3.2: Latency and throughputs as measures by the middlewares

5.2.1 Explanation

Provide a detailed analysis of the results (e.g., bottleneck analysis, component utilizations, average queue lengths, system saturation). Add any additional figures and experiments that help you illustrate your point and support your claims.

5.3 Histogram

For the case with 6 keys inside the multi-get, display four histograms representing the sharded and non-sharded response time distribution, both as measured on the client, and inside the middleware. Choose the bucket size in the same way for all four, and such that there are at least 10 buckets on each of the graphs.

5.4 Summary

Provide a detailed comparison of the sharded and non-sharded modes. For which multi-GET size is sharding the preferred option? Provide a detailed analysis of your system. Add any additional figures and experiments that help you illustrate your point and support your claims.

6 2K Analysis (90 pts)

This 2k analysis includes the analysis of any correlations between the following factors: 1. number of memcached servers, 2. the number of middleware vm's and finally 3. the number of worker threads per middleware. The following table shows which possible values we cross-reference, such that we can later on analyse which factors have the most impact on throughput and response time.

- Memcached servers: 1 and 3
- Middlewares: 1 and 2
- Worker threads per MW: 8 and 32

Any configuration is run 3 times (3 repetitions) for 90 seconds, which implies a 15 second warm-up and a 15 second cool-down time. The following table shows a detailed configuration of my setup.

Number of servers	1 and 3
Number of client machines	3
Instances of memtier per machine	1 (1 middleware) or 2 (2 middlewares)
Threads per memtier instance	2 (1 middleware) or 1 (2 middlewares)
Virtual clients per thread	32
Workload	Write-only and Read-only
Number of middlewares	1 and 2
Worker threads per middleware	8 and 32
Repetitions	3 or more (at least 1 minute each)

I apply this analysis once for read-only workloads, and then separately for write-only workloads, as both procedures are fundamentally different. I don't use any multi-get behavior (i.e. keysize is always 1 for read-only workloads).

For read-only and work-only workloads, I created tables which represent all possible configurations. For both subsections, I will use the following abbreviations (and variable-names)

- NS: Number of Servers
- NM: Number of Middlewares
- WTMW: Worker threads per middleware

6.0.1 Read-only

NS (Servers)	NM (Middlewares)	WTMW (Workerthreads)	Mean Throughput
1	1	8	
1	1	32	
1	2	8	
1	2	32	
3	1	8	
3	1	32	
3	2	8	
3	2	32	

Repeat the experiment for (a) a write-only and (b) a read-only workload. For each of the two workloads, what is the impact of these parameters on throughput, respectively response time?

6.0.2 Write-only

7 Queuing Model (90 pts)

In this section I model the workerqueue of the middleware (after the requests come in) using queueing theory to model how the system behaves with an asymptotically increasing number of threads. In both subsection I will go use the different number of middleware (specifically, one of [8, 16, 32, 64]) threads to apply this analysis.

7.1 M/M/1

In this subsection I model the behavior of the middleware and it's workerqueue using a M/M/1 queueing model.

I choose the following input parameters to model the system:

First of all, I define the service rate to the model. For this, I look at the maximum throughput per number of middleware threads (amongst all given number of virtual clients). From the file `create_lineplot_exp4_1.py` I read out the respective maximum throughput values:

Threads in the middleware	Maximum Throughput (ops/sec)
8	14365.88
16	18338.97
32	18828.88
64	17984.48

your entire system. Motivate your choice of input parameters to the model. Explain for which experiments the predictions of the model match and for which they do not.

7.2 M/M/m

Build an M/M/m model based on Section 4, where each middleware worker thread is represented as one service. Motivate your choice of input parameters to the model. Explain for which experiments the predictions of the model match and for which they do not.

7.3 Network of Queues

Based on Section 3, build a network of queues which simulates your system. Motivate the design of your network of queues and relate it wherever possible to a component of your system. Motivate your choice of input parameters for the different queues inside the network. Perform a detailed analysis of the utilization of each component and clearly state what the bottleneck of your system is. Explain for which experiments the predictions of the model match and for which they do not.